

# **Advanced Algorithms**

**Course Code: 21CS7E12**  
**Module 4**

## **TEXT BOOK**

**Introduction to Algorithms- Thomas H. Cormen, Charles  
E. Leiserson, Ronald L. Rivest, Clifford Stein PHI, 3<sup>rd</sup>  
Edition, 2009**

# **Module 4**

## **B-trees and Fibonacci Heaps**

- 1.Introduction and Definition of B-trees
- 2 Basic operations of B-trees
- 3 Introduction to Fibonacci heaps
- 4 Structure of Fibonacci heaps
- 5 Mergeable-heap operations

# Module 4


## B Tree

- 1.Introduction and Definition of B-trees
- 2 Basic operations of B-trees

# Introduction to B-Tree


- B-trees are balanced search tree.
- More than two children are possible.
- B-Tree, stores all information in the leaves and stores only keys and Child pointer.
- If an internal B-tree node  $x$  contains  $n[x]$  keys then  $x$  has  $n[x]+1$  children.

# Example of B-Tree




**Figure 18.1** A B-tree whose keys are the consonants of English. An internal node  $x$  containing  $n[x]$  keys has  $n[x] + 1$  children. All leaves are at the same depth in the tree. The lightly shaded nodes are examined in a search for the letter  $R$ .

# Another example of B-Tree



**Figure 18.3** A B-tree of height 2 containing over one billion keys. Each internal node and leaf contains 1000 keys. There are 1001 nodes at depth 1 and over one million leaves at depth 2. Shown inside each node  $x$  is  $n[x]$ , the number of keys in  $x$ .

# Application of B-Tree



It designed to work on magnetic disks or other (direct access) secondary storage devices.

# Properties of B-Tree

- A B-Tree  $T$  is a rooted tree having the following properties:
  1. Every node  $x$  has the following fields:
    1.  $n[x]$  the no. of keys currently stored in node  $x$
    2. The  $n[x]$  keys themselves, stored in non-decreasing order, so that  $key_1[x] \leq key_2[x] \dots \leq key_{n-1}[x] \leq key_n[x]$ .
    3.  $Leaf[x]$ , a Boolean value that is TRUE if  $x$  is a leaf and FALSE if  $x$  is an internal node.
  2. Each internal node  $x$  also contains  $n[x]+1$  pointers (Childs)  $c_1[x], c_2[x], \dots, c_{n[x]+1}[x]$ .
  3. All leaves have the same depth, which is the tree's height  $h$ .



# Properties of B-Tree (cont.)

4. There are lower and upper bounds on the no. of keys a node can contains: these bounds can be expressed in terms of a fixed integer  $t \geq 2$  called the minimum degree of B-Tree.
- Every node other than the root must have at least  $t-1$  keys, then root has at least  $t$  children if the tree is non empty the root must have at least one key.
  - Every node can contain at most  $2t-1$  keys. Therefore, an internal node can have at most  $2t$  children we say that a node is full if it contains exactly  $2t-1$  keys.

# Height of B-tree


- Theorem:

If  $n \geq 1$ , then for any  $n$ -key B-tree  $T$  of height  $h$  and minimum degree  $t \geq 2$ ,

$$h \leq \log_t (n+1)/2$$

- Proof:

- The root contains at least one key
- All other nodes contain at least  $t-1$  keys.
- There are at least 2 nodes at depth 1, at least  $2t$  nodes at depth 2, at least  $2t^{i-1}$  nodes at depth  $i$  and  $2t^{h-1}$  nodes at depth  $h$



**Figure 18.4** A B-tree of height 3 containing a minimum possible number of keys. Shown inside each node  $x$  is  $n[x]$ .

# Basic operation on B-tree

- **B-TREE-SEARCH** :-Searching in B Tree
- **B-TREE-INSERT** :-Inserting key in B Tree
- **B-TREE-CREATE** :-Creating a B Tree
- **B-TREE-DELETE** :- Deleting a key from B Tree

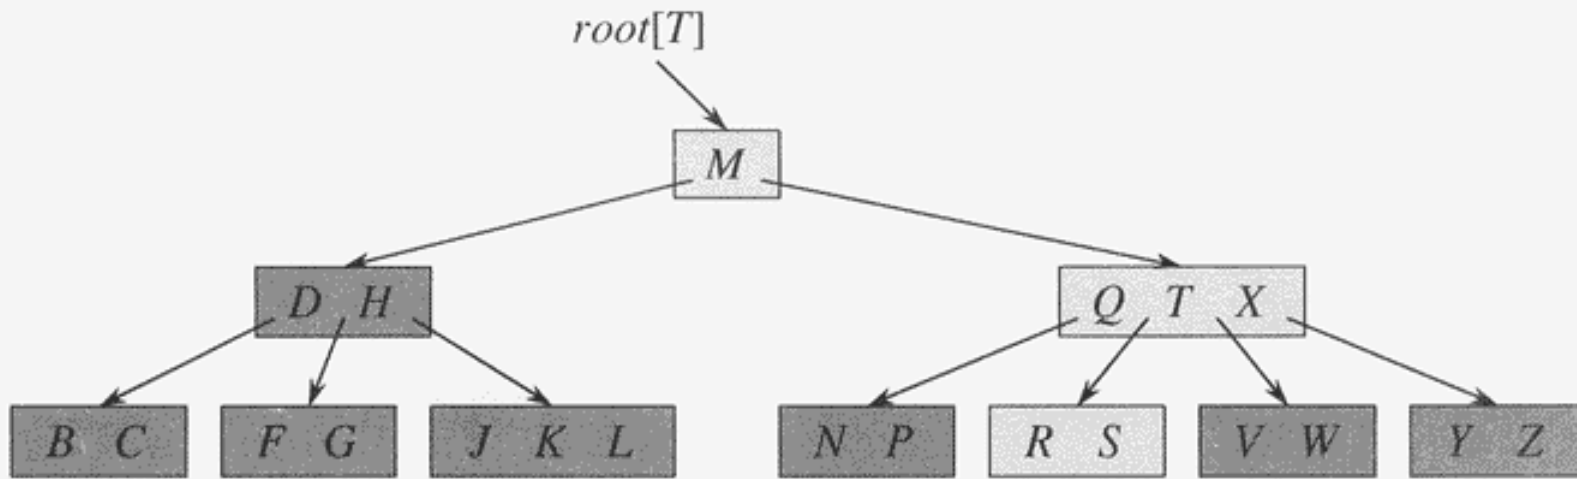
# X is a subtree and k is searching element

B-TREE-SEARCH( $x, k$ )

Complexity= $O(t)$

```
1   $i \leftarrow 1$ 
2  while  $i \leq n[x]$  and  $k > key_i[x]$ 
3      do  $i \leftarrow i + 1$ 
4  if  $i \leq n[x]$  and  $k = key_i[x]$ 
5      then return  $(x, i)$ 
6  if  $leaf[x]$ 
7      then return NIL
8  else DISK-READ( $c_i[x]$ )
9      return B-TREE-SEARCH( $c_i[x], k$ )
```

Complexity=  
 $O(\log_t n)$



# Creating an empty B tree


**B-TREE-CREATE( $T$ )**

```
1   $x \leftarrow \text{ALLOCATE-NODE}()$   
2   $leaf[x] \leftarrow \text{TRUE}$   
3   $n[x] \leftarrow 0$   
4   $\text{DISK-WRITE}(x)$   
5   $root[T] \leftarrow x$ 
```

# Insert

- Cannot just create a new leaf node and insert it
  - resulting tree is not B-tree
- Insert new key into an existing leaf node
- If leaf node is full
  - Split full node  $y$  (with  $2t-1$ ) keys around its median  $\text{key}_t[y]$  into two nodes each having  $t-1$  keys
  - Move the median key into  $y$ 's parent.
  - If parent is full, recursively split, all the way to the root node if necessary.
  - If root is full, split root - height of tree increase by one.


# INSERT



**Figure 18.1** A B-tree whose keys are the consonants of English. An internal node  $x$  containing  $n[x]$  keys has  $n[x] + 1$  children. All leaves are at the same depth in the tree. The lightly shaded nodes are examined in a search for the letter  $R$ .



# Split



- Splitting a node with  $t=4$

# B-TREE-SPLIT-CHILD ALGORITHM

B-TREE-SPLIT-CHILD( $x, i, y$ )

1.  $z \leftarrow \text{ALLOCATE-NODE}()$
2.  $\text{leaf}[z] \leftarrow \text{leaf}[y]$
3.  $n[z] \leftarrow t-1$
4. for  $j \leftarrow 1$  to  $t-1$
5.     do  $\text{key}_j[z] \leftarrow \text{key}_{j+t}[y]$
6. if not  $\text{leaf}[y]$
7.     then for  $j \leftarrow 1$  to  $t$
8.         do  $c_j[z] \leftarrow c_{j+t}[y]$
9.  $n[y] \leftarrow t-1$

# cont.....

10. for  $j \leftarrow n[x] + 1$  downto  $i + 1$
11.     do  $c_{j+1}[x] \leftarrow c_j[x]$
12.  $c_{j+1}[x] \leftarrow z$
13. for  $j \leftarrow n[x]$  downto  $i$
14.     do  $key_{j+1}[x] \leftarrow key_j[x]$
15.  $key_i[x] \leftarrow key_t[y]$
16.  $n[x] \leftarrow n[x] + 1$
17. DISK-WRITE( $y$ )
18. DISK-WRITE( $z$ )
19. DISK-WRITE( $x$ )

# B-TREE-INSERT ALGORITHM

B-TREE-INSERT( $T, k$ )

1.  $r \leftarrow \text{root}[T]$
2. If (  $n[r] = 2t-1$ )
3.     **then**  $s \leftarrow \text{Allocate-Node}()$
4.          $\text{root}[T] \leftarrow s$
5.          $\text{leaf}[s] \leftarrow \text{FALSE}$
6.          $n[s] \leftarrow 0$
7.          $c_1[s] \leftarrow r$
8.         B-TREE-SPLIT-CHILD( $s, 1, r$ )
9.         B-TREE-INSERT-NONFULL( $s, k$ )
10.     **else** B-TREE-INSERT-NONFULL( $r, k$ )

## **Algorithm: Insert Key into B-Tree**

### **1. Start at the Root:**

Begin at the root node of the B-tree.

### **2. Traverse the Tree:**

- Find the appropriate leaf node where the new key should be inserted.
- While traversing:
  - If the current node is a leaf, stop.
  - If the current node is not a leaf, follow the appropriate child pointer based on the value of the key compared to the keys in the current node.

### **3. Check Space in the Leaf Node:**

- If the leaf node has fewer than  $2t-1$  keys:
  - Insert the key in sorted order in the node.
- Otherwise:
  - The leaf node is full. Split it and propagate a key to the parent.

#### **4. Splitting a Full Node:**

- If a node overflows (contains  $2t$  keys), split it:
  - Divide the  $2t$  keys into two nodes:
    - Left node: Contains the first  $t-1$  keys.
    - Right node: Contains the last  $t-1$  keys.
  - Promote the middle key (the  $t$ -th key) to the parent node.
- If the parent node also overflows due to this promotion, repeat the splitting process upward, possibly splitting the root.

#### **5. Handle Root Overflow:**

- If the root node overflows, create a new root:
  - The new root contains the promoted middle key.
  - The two split nodes become its children.
- This increases the height of the B-tree by 1.

#### **6. Repeat Until Key is Inserted:**

- Continue the process until the key is

# Process of finding the appropriate position for the key in the tree.

## 1. Start from the Root:

- Compare the key with the keys in the current node.
- Traverse the child pointer corresponding to the key range where the new key belongs.

## 2. Repeat Until a Leaf is Reached:

- For each internal node encountered:
  - Find the correct child node by comparing the key with the keys in the node.
  - Move to that child node.

## 3. Insert into the Leaf Node:

- Once at the appropriate leaf, insert the key into the correct position based on sorted order.

## Example

Consider inserting the key 15 into a B-tree with  $t=3$ .

Traverse from the root to the correct child node where 15 belongs.


If the leaf node is not full, insert 15 in sorted order.

If the leaf node is full:

Split the node, promoting the middle key to the parent.

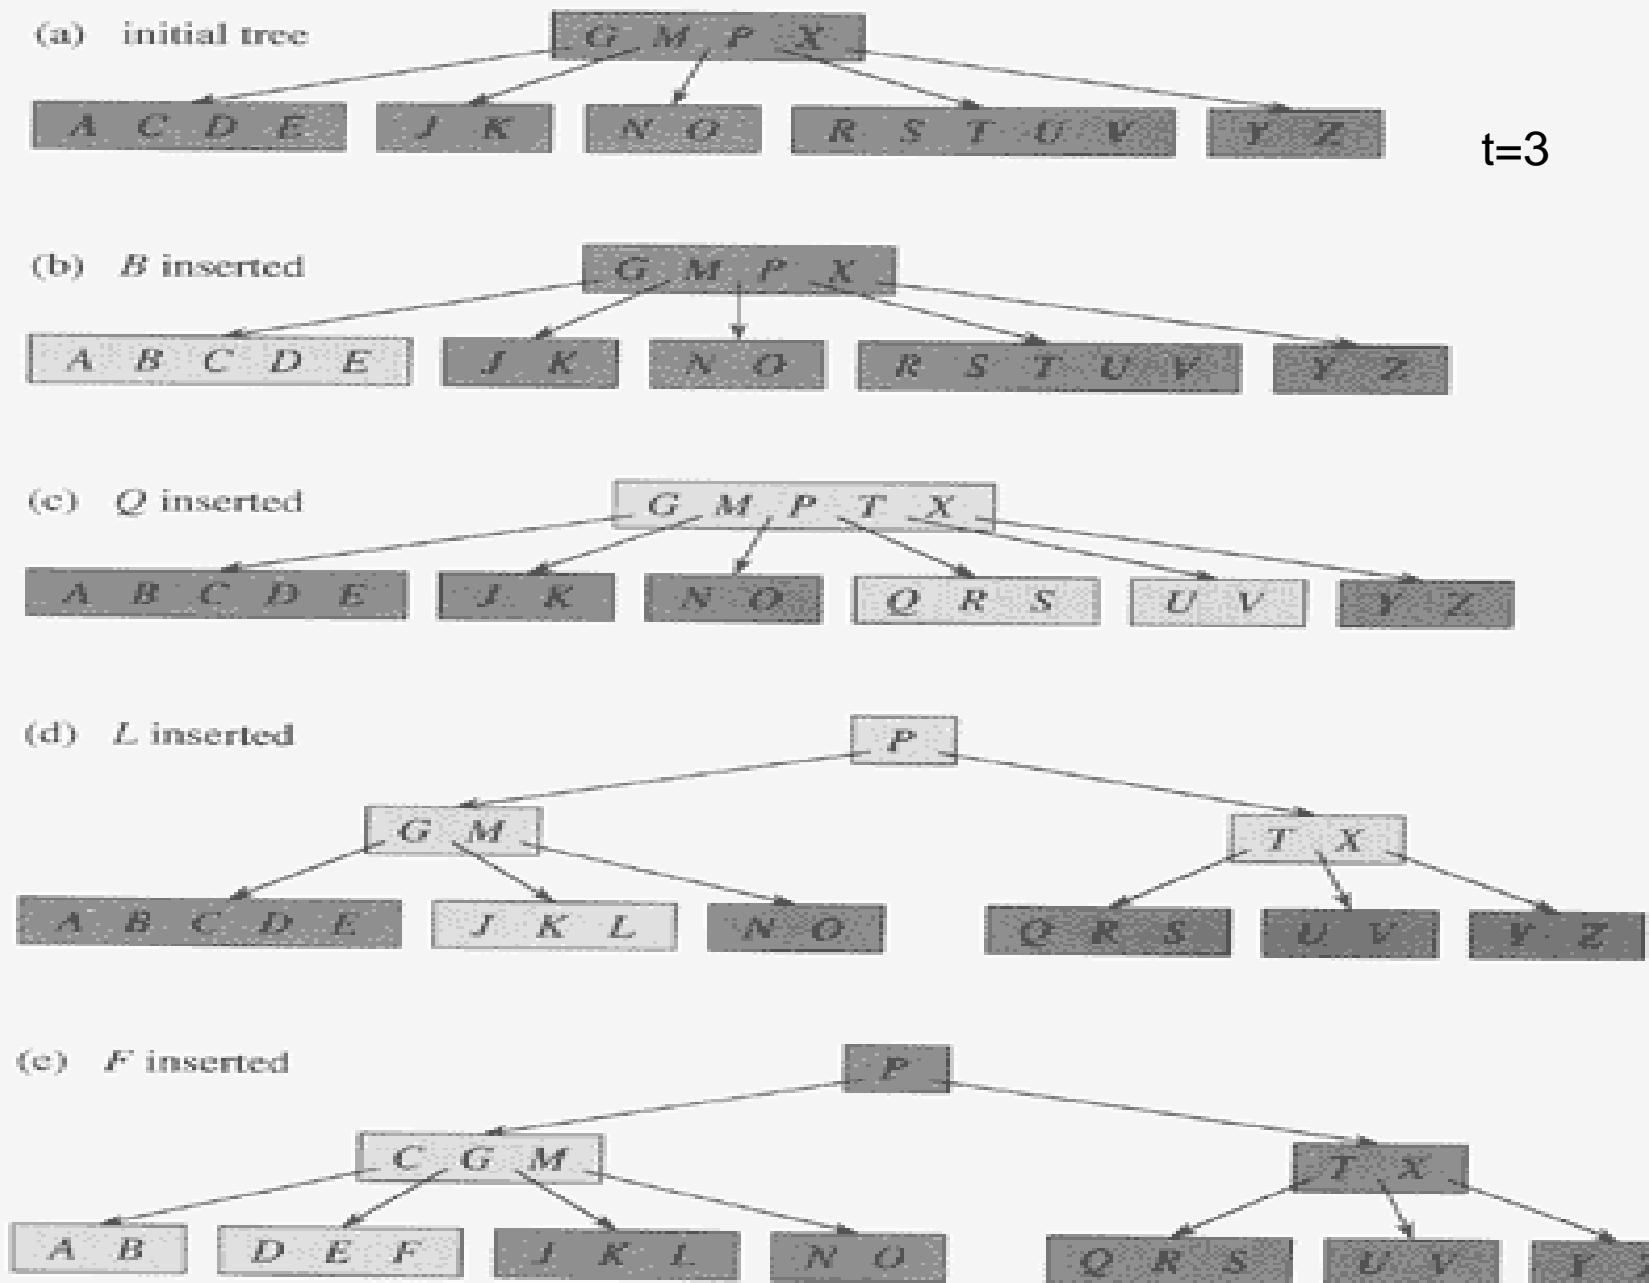
Adjust pointers and repeat if necessary.

- Can also insert in a single pass down the tree, instead of going down during search then recursively splitting on the way up
  - Split full nodes encountered on the way down during search.
  - Example:



Initial tree,  $t=3$

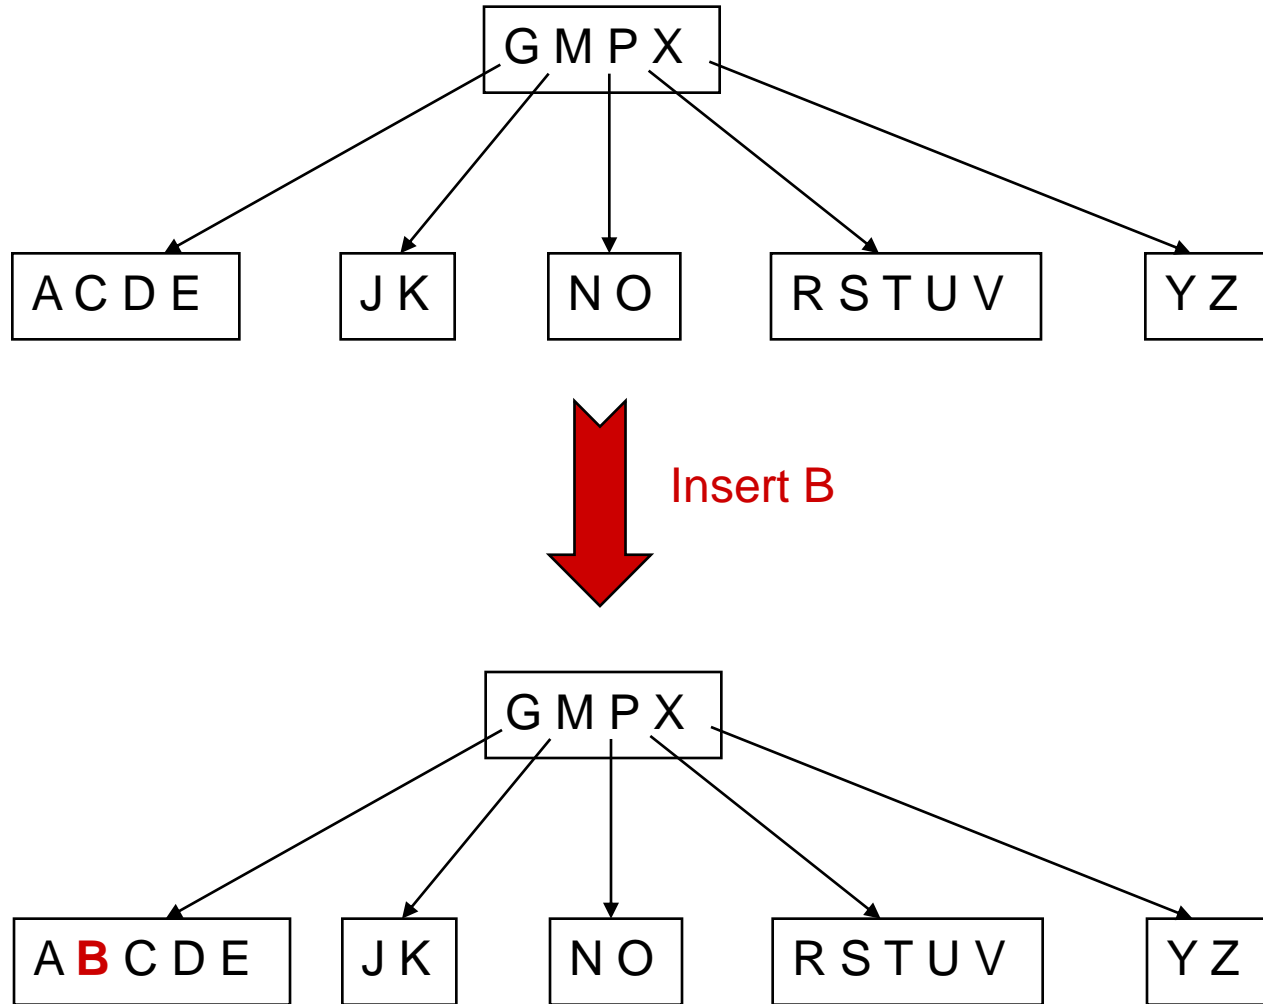




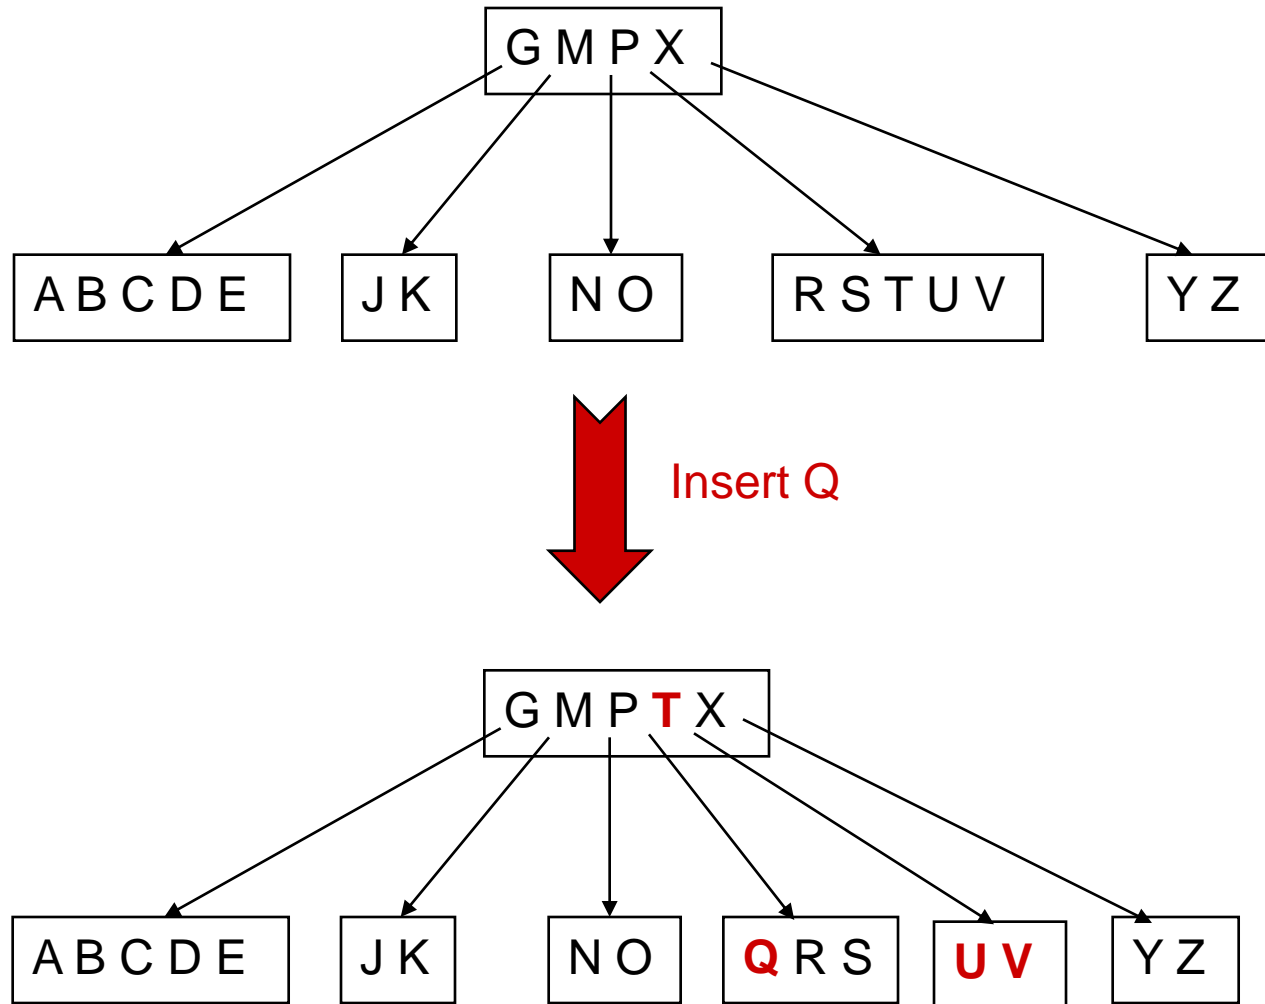
**Figure 18.7** Inserting keys into a B-tree. The minimum degree  $t$  for this B-tree is 2, because the

# Insert Example

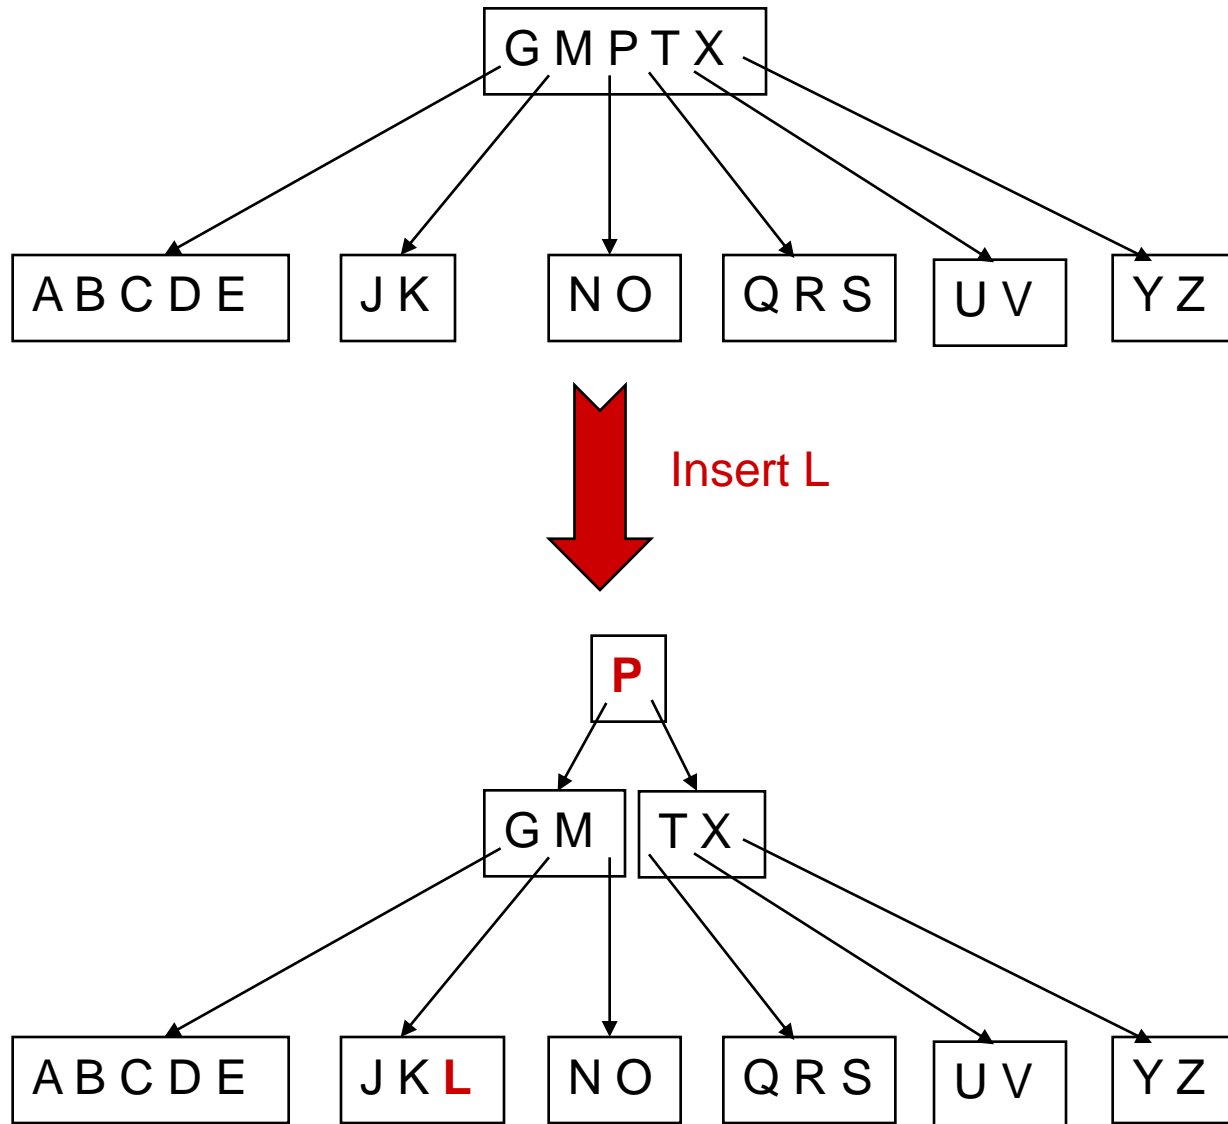
t = 3



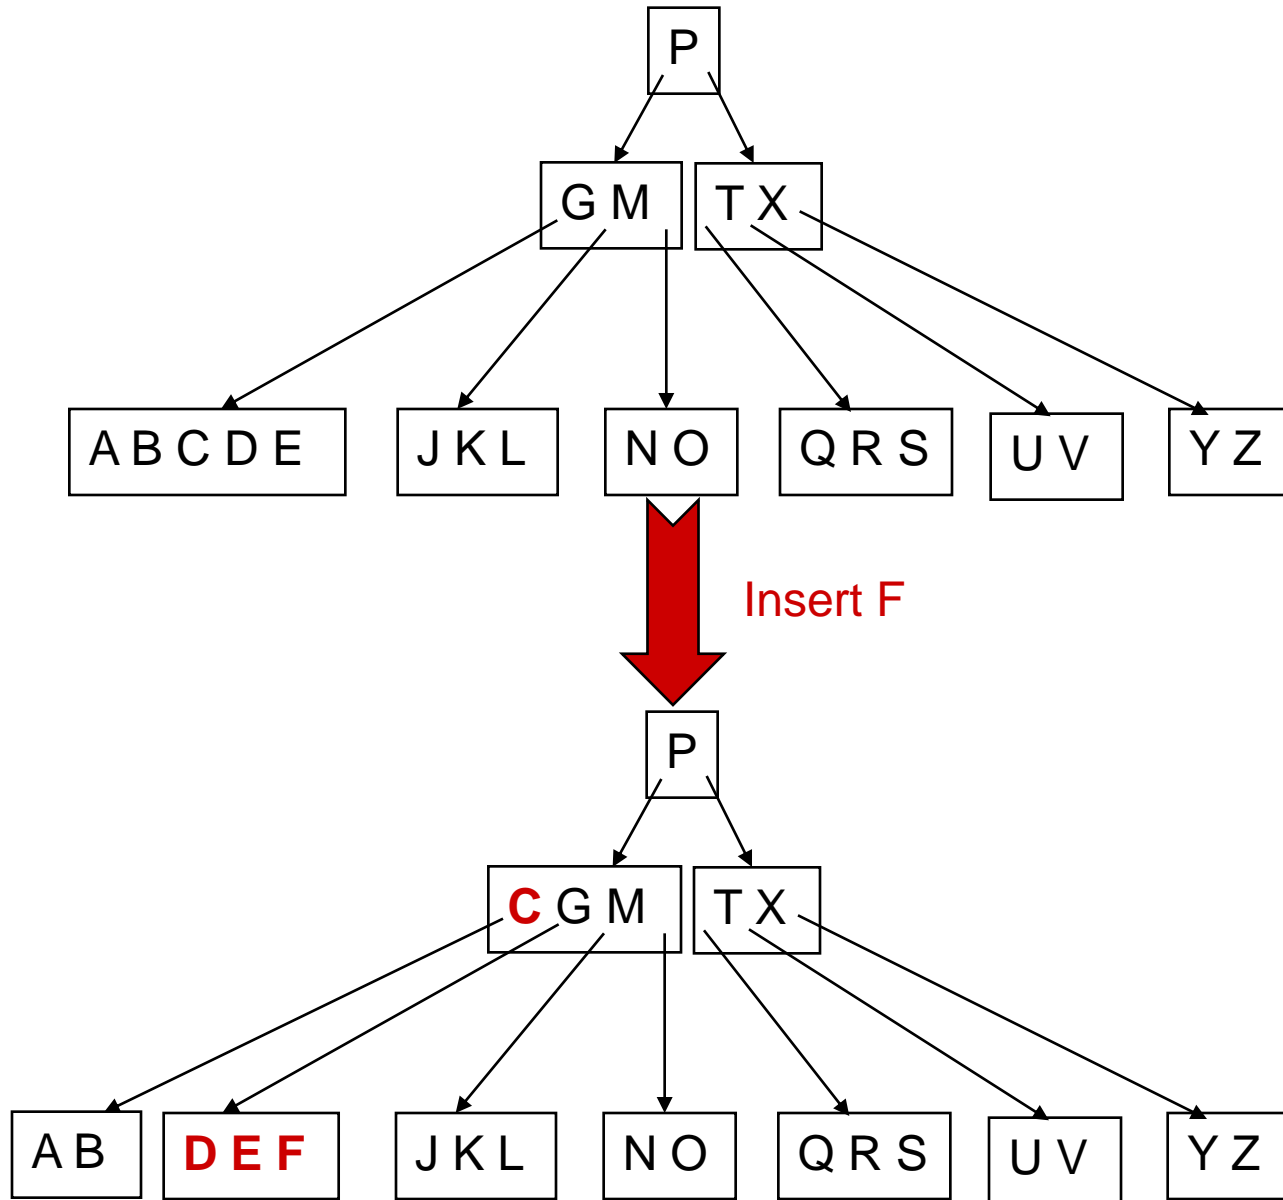
# Insert Example (Continued)



# Insert Example (Continued)




# Insert Example (Continued)



# Constructing a B-tree


- Suppose we start with an empty B-tree and keys arrive in the following order: 1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55 45
- We want to construct a B-tree of degree 3
- The first four items go into the root:



- To put the fifth item in the root would violate condition 5
- Therefore, when 25 arrives, pick the middle key to make a new root

# Constructing a B-tree

Add 6 to the tree




Exceeds Order.  
Promote middle and  
split.


1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)



6, 14, 28 get added to the leaf nodes:






1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)


Adding 17 to the right leaf node would over-fill it, so we take the middle key, promote it (to the root) and split the leaf



1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)


7, 52, 16, 48 get added to the leaf nodes



1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)


Adding 68 causes us to split the right most leaf, promoting 48 to the root



1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)


Adding 3 causes us to split the left most leaf



1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)

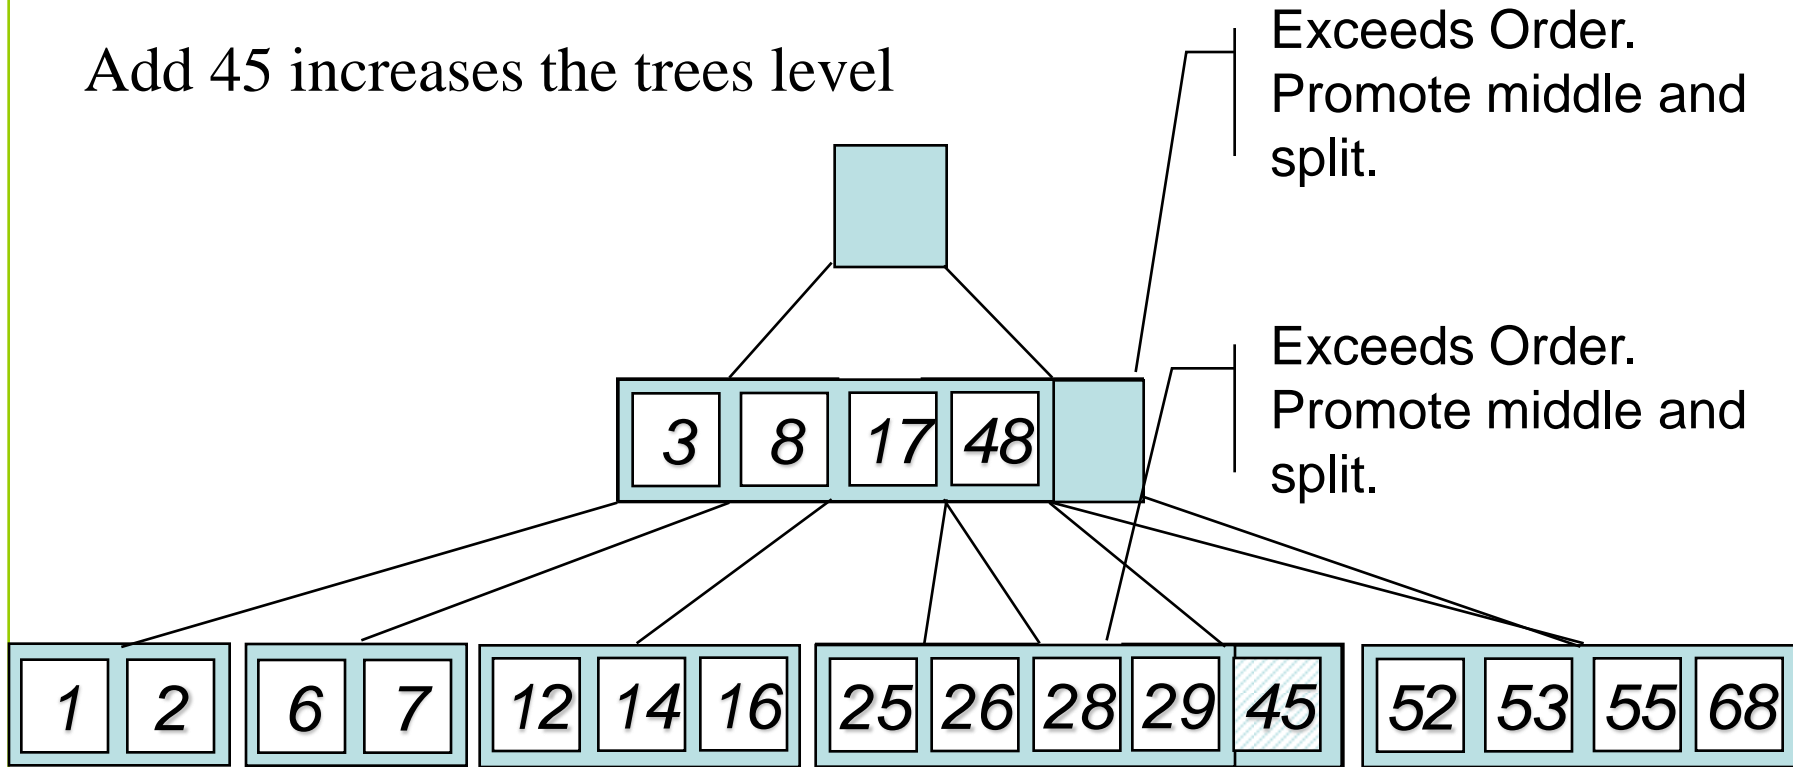
Add 26, 29, 53, 55 then go into the leaves



1  
12  
8  
2  
25  
6  
14  
28  
17  
7  
52  
16  
48  
68  
3  
26  
29  
53  
55  
45

# Constructing a B-tree (contd.)

Add 45 increases the trees level



# Exercise in Inserting a B-Tree

1) Insert the following keys in B-tree when  $t=3$  :

- 3, 7, 9, 23, 45, 1, 5, 14, 25, 24, 13, 11, 8, 19, 4, 31, 35, 56

2) Difference between B-Tree and Red black tree.

## Aspect

## Red-Black Tree

## B-Tree

### Definition

A type of self-balancing binary search tree.

A self-balancing search tree designed for databases and storage systems.

### Node Structure

Each node has a single key (binary search tree).

Each node can store multiple keys (min:  $t-1$ , max:  $2t-1$ ).

### Child Structure

Each node has at most two children.

Each node can have multiple children (min:  $t$ , max:  $2t$ ).

### Balancing Method

Balance is maintained using **coloring** (red/black) and **rotations** during insertions and deletions.

Balance is maintained by splitting or merging nodes during insertions and deletions.

### Height

Height is approximately  $O(\log_{\frac{n}{t}} n)$  (binary tree).

Height is also  $O(\log_{\frac{n}{t}} n)$ , but generally shorter than a red-black tree due to larger branching factor.



## Aspect

### Use Case

### Traversal Complexity

### Search Efficiency

### Space Utilization

### Memory Access

### Rotations

### Applications

## Red-Black Tree

Used for in-memory data structures (e.g., compilers, network routing).

Traversal involves multiple nodes (binary structure).

Search time is  $O(\log n)$ , but nodes are smaller.

Stores one key per node, less space-efficient.

Frequently accesses small amounts of data, which is CPU-cache-friendly.

Requires frequent rotations during balancing.

- Symbol tables in compilers.
- Network routing tables.

## B-Tree

Used for disk-based storage systems (e.g., databases, filesystems).

Traversal is efficient due to fewer nodes and larger fanout.

Search time is  $O(\log \frac{n}{b})$ , with fewer disk I/O operations due to multi-key nodes.

Optimized for storage efficiency with multiple keys per node.

Designed for minimizing disk I/O with fewer accesses and larger blocks.

No rotations; rebalancing involves splits or merges.

- Database indexing (e.g., MySQL, PostgreSQL).
- Filesystem indexing (e.g., NTFS, EXT).

## Example Problem

**As a function of the minimum degree  $t$ , what is the maximum number of keys that can be stored in a B-tree of height  $h$ ?**

**Key properties of a B-tree:**

**Minimum degree  $t$ :** This is the minimum number of children a node in the B-tree can have.

Every internal node (except the root) has at least  $t$  children.

Each internal node can have at most  $2t-1$  keys and  $2t$  children.

**Height  $h$ :** This is the number of levels in the tree, with the root being at height 0.

**At level  $i$ ,** there can be up to  $(2t)^i$  nodes, and each of these nodes can hold up to  $2t-1$  keys.

**Total number of keys:**

**At level 0,** the root can hold up to  $2t-1$  keys.

**At level 1,** there are up to  $2t$  nodes, each holding up to  $2t-1$  keys.

**At level 2,** there are up to  $(2t)^2$  nodes, each holding up to  $2t-1$  keys.

This continues until **level h**,

where Number of nodes =  $1 + 2t + (2t)^2 + \dots + (2t)^h$

Here,

First term (a) = 1

Common ratio (r) = 2t

Number of terms = h+1

Sum of the Geometric Series  
is given by:

$$S = a \cdot (r^{n+1} - 1) / (r - 1)$$

Substituting these values:

$$S = ((2t)^{h+1} - 1) / (2t - 1)$$

### **Maximum Number of Keys**

Each node can have at most 2t-1  
keys.

So, the maximum number of keys

**N<sub>max</sub>** is:

$$N_{\max} = (S - 1) \cdot (2t - 1)$$

After simplifying:

$$N_{\max} = (2t)^{h+1} - 1$$

This represents the maximum  
number of keys a B-tree of height  
h and minimum degree t can hold.

## Example Problem

**Suppose that we insert the keys  $\{1,2,\dots,n\}$  into an empty B-tree with minimum degree 2. How many nodes does the final B-tree have?**

### B-tree Structure

#### Minimum Degree $t=2$ :

Each node (except the root) contains at least  $t-1$  key and at most  $2t-1=3$  keys.

The root can have fewer keys but must have at least 1 key (if not empty).

Each internal node has between 2 and 4 children (one more than the number of keys).

#### Node Splitting:

- A node splits when it becomes full, i.e., when it contains  $2t-1=3$  keys.
- During splitting:
  - The middle key moves up to the parent.
  - The node splits into two child nodes.
- This increases the number of nodes in the tree.

Contd.....

## Step-by-Step Analysis

### 1. Keys in the Root

- Initially, the root node can hold up to  $2t-1=3$  keys.
- When the root splits, it creates two child nodes, and the tree grows in height.

### 2. Number of Splits

- For  $n$  keys, the total number of splits determines the number of nodes.
- Each split adds a new node to the tree.
- At any point in the tree's growth:
  - The number of nodes in the tree is related to how keys are distributed among all levels.

### 3. Approximation for Total Nodes

- When inserting  $n$  keys into a B-tree:
  - A key distributes across nodes to ensure balance.
  - Each node can hold approximately **2 keys** on average (halfway between  $t-1$  and  $2t-1$ ).
- The total number of nodes is approximately:  $\text{Number of Nodes} \approx \lceil n/2 \rceil$

# B-TREE-INSERT-NONFULL( $x, k$ )

```
1   $i \leftarrow n[x]$ 
2  if  $leaf[x]$ 
3      then while  $i \geq 1$  and  $k < key_i[x]$ 
4          do  $key_{i+1}[x] \leftarrow key_i[x]$ 
5               $i \leftarrow i - 1$ 
6           $key_{i+1}[x] \leftarrow k$ 
7           $n[x] \leftarrow n[x] + 1$ 
8          DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < key_i[x]$ 
10      do  $i \leftarrow i - 1$ 
11       $i \leftarrow i + 1$ 
12      DISK-READ( $c_i[x]$ )
13      if  $n[c_i[x]] = 2t - 1$ 
14          then B-TREE-SPLIT-CHILD( $x, i, c_i[x]$ )
15              if  $k > key_i[x]$ 
16                  then  $i \leftarrow i + 1$ 
17      B-TREE-INSERT-NONFULL( $c_i[x], k$ )
```

# Deleting from B-Trees

# The Concept


- You can delete a key entry from any node.
- ->Therefore, you must ensure that before/after deletion, the B-Tree maintains its properties.
- When deleting, you have to ensure that a node doesn't get *too small* (minimum node size is  $T - 1$ ). We prevent this by *combining* nodes together.



# Lets look at an example:




We're given this valid B-Tree

Note:  $T = 3$



# Deletion Cases


- **Case 1:** If the key  $k$  is in node  $x$  and  $x$  is a leaf node having at least  $t$  keys - then delete  $k$  from  $x$ .



# Simple Deletion


Case 1: We delete “F”

Result: We remove “F” from the leaf node. No further action needed.




# Deletion Cases (Continued)

- **Case 2:** If the child key  $k$  is in node  $x$  and  $x$  is an internal node, do the following:



# Deletion Cases (Continued)

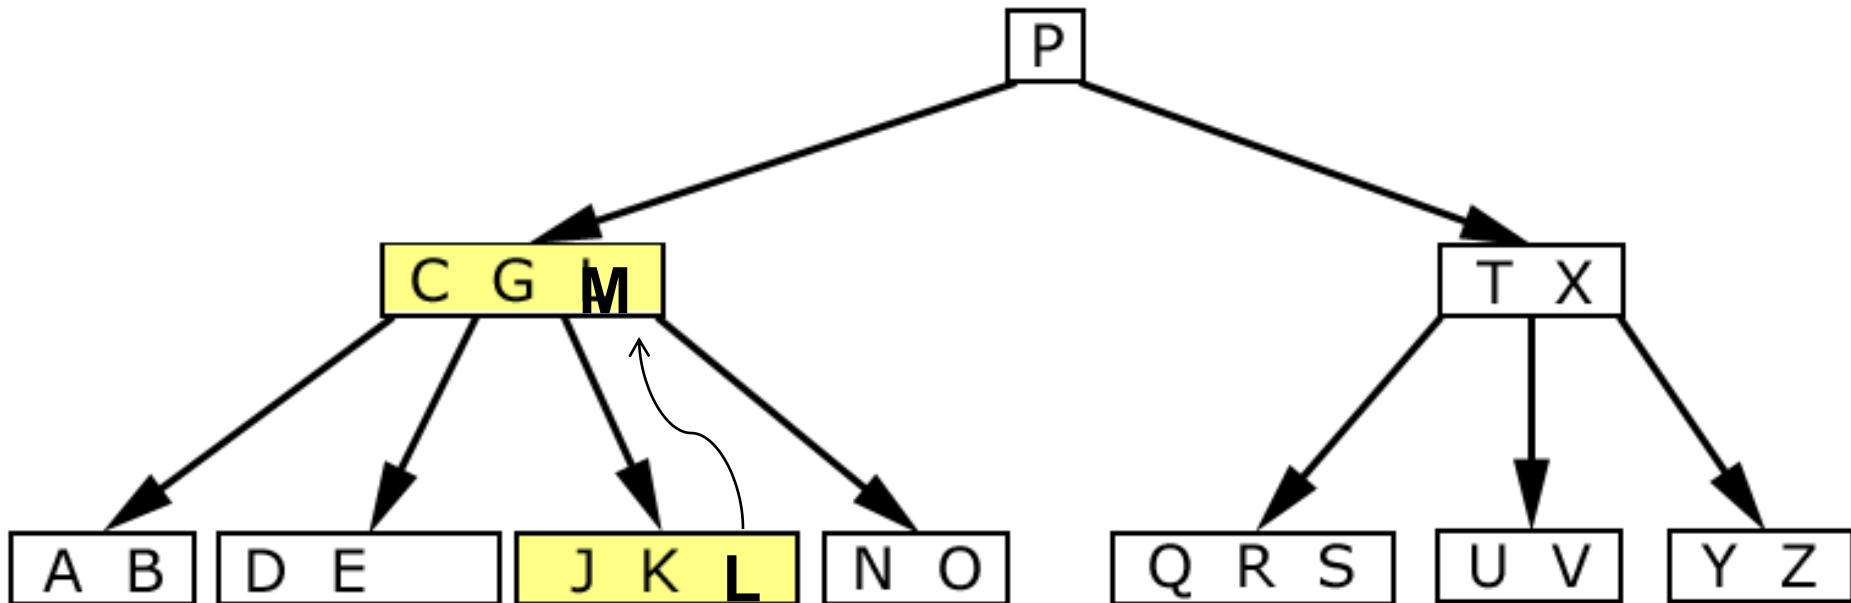
- **Subcase a:** If the child  $y$  that precedes  $k$  has at least  $t$  keys then find predecessor  $k'$  of  $k$  in subtree rooted at  $y$ , recursively delete  $k'$  and replace  $k$  by  $k'$  in  $x$ .



# Deleting and shifting


Case 2a: We deleted “M”

Result: We remove “M” from the parent node. Since there are four nodes and two letters, we move “L” to replace “M”. Now, the “N O” node has a parent again.




# Deletion Cases (Continued)

**Subcase B:** Symmetrically, if the child  $z$  that follows  $k$  in node  $x$  has at least  $t$  keys then find successor  $k'$  of  $k$  in subtree rooted at  $z$ , recursively delete  $k'$  and replace  $k$  by  $k'$  in  $x$ .



# Deletion Cases (Continued)

**Subcase C:**  $y$  and  $z$  both have  $t-1$  keys -- merge  $k$  and  $z$  into  $y$ , free  $z$ , recursively delete  $k$  from  $y$ .






# Combining and Deleting

Case 2c: Now, we delete “G”

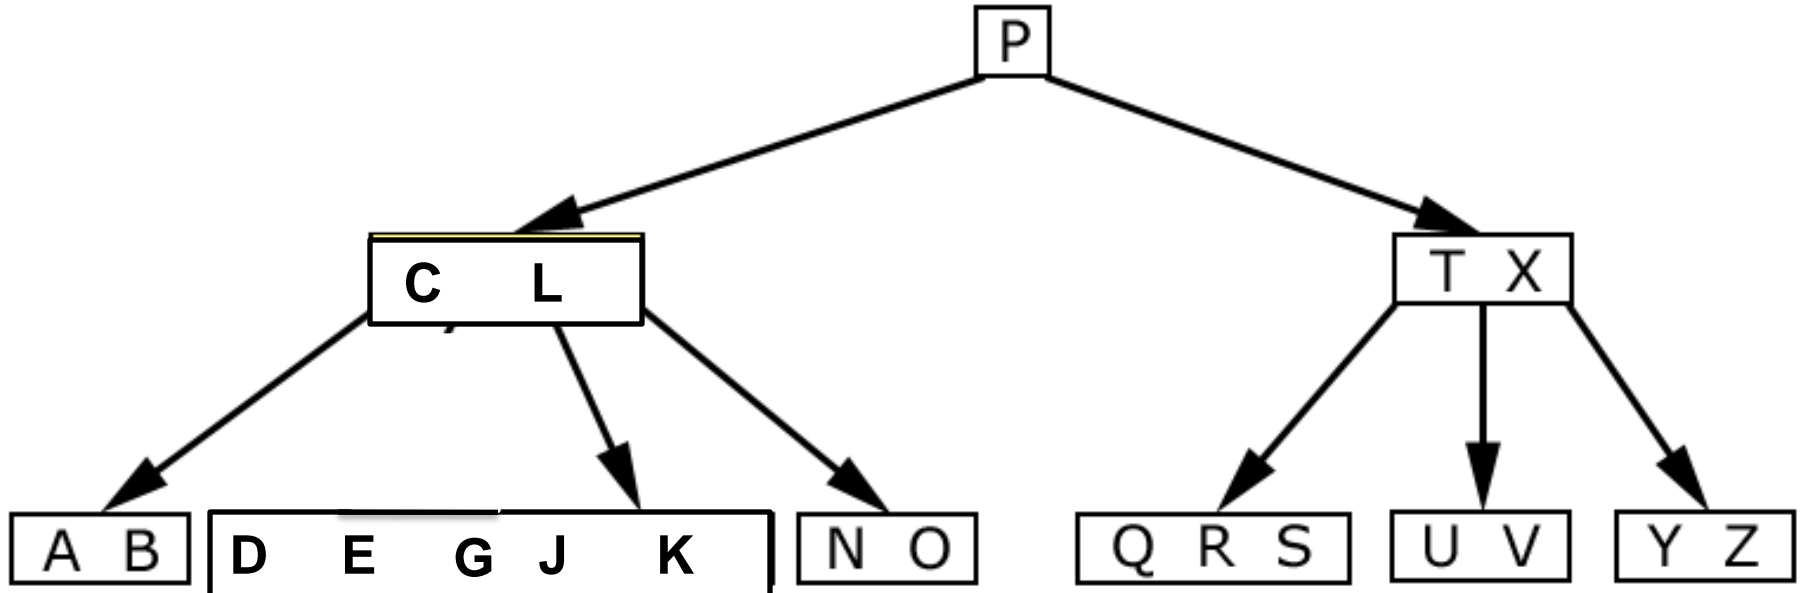
Result: First, we combine nodes “DE” and “JK”. Then, we push down “G” into the “DEJK” node and delete it as a leaf.



# Combining and Deleting

Case 2c: Now, we delete “G”

Result: First, we combine nodes “DE” and “JK”. Then, we push down “G” into the “DEJK” node and delete it as a leaf.




# Deletion Cases (Continued)

**Case 3:**  $k$  not in internal node. Let  $c_i[x]$  be the root of the subtree that must contain  $k$ , if  $k$  is in the tree. If  $c_i[x]$  has at least  $t$  keys, then recursively descend; otherwise, execute 3.A and 3.B as necessary.


# Deletion Cases (Continued)

**Subcase A:**  $c_i[x]$  has  $t-1$  keys, some sibling has at least  $t$  keys.




# Deleting “B”

Before:




After: Deleted “B”, Demoted “C”, Promoted “E”



# Deletion Cases (Continued)

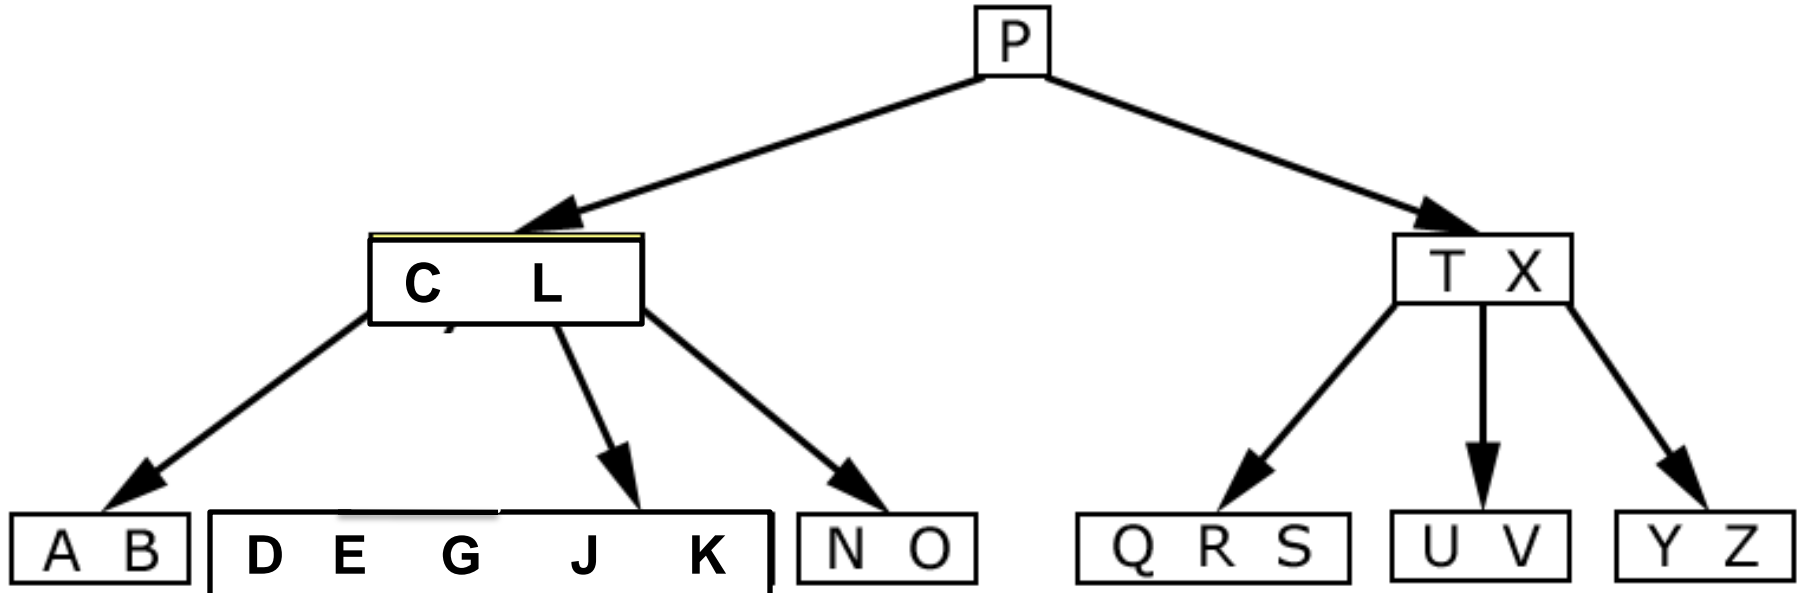
**Subcase B:**  $c_i[x]$  and sibling both have  $t-1$  keys.



# Combining and Deleting


Case 3b: Now, we delete “D”

Result: First, we combine nodes “DE” and “JK”. Then, we push down “G” into the “DEJK” node and delete “D” as a leaf.



# Type #1: Simple leaf deletion


Assuming a 5-way  
B-Tree, as before...




Delete 2: Since there are enough  
keys in the node, just delete it




# Type #2: Simple non-leaf deletion




# Type #4: Too few keys in node and its siblings



# Type #4: Too few keys in node and its siblings



# Type #3: Enough siblings



# Type #3: Enough siblings

